

ТМО ЛР5 ИУ5-63Б Горкунов Николай

5 июня 2024 г.

1 ТМО ЛР5 ИУ5-63Б Горкунов Николай

2 Ансамбли моделей машинного обучения. Часть 1.

- Выберите набор данных (датасет) для решения задачи классификации или регрессии.
- В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.
- С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.
- Обучите следующие ансамблевые модели:
 - две модели группы бэггинга (бэггинг или случайный лес или сверхслучайные деревья);
 - AdaBoost;
 - градиентный бустинг.
- Оцените качество моделей с помощью одной из подходящих для задачи метрик. Сравните качество полученных моделей.

3 Набор данных: Boston housing dataset

```
[1]: import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
from io import StringIO
from PIL import Image
from IPython.display import display
import graphviz
import pydotplus
from sklearn.tree import DecisionTreeRegressor, export_graphviz
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import ExtraTreesRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.ensemble import BaggingRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.pipeline import make_pipeline
```

```
import seaborn as sns
import time
import matplotlib.pyplot as plt
from kaggle.api.kaggle_api_extended import KaggleApi
pd.options.display.max_columns = None
```

```
[2]: kaggle_api = KaggleApi()
kaggle_api.authenticate()
kaggle_api.dataset_download_files('altavish/boston-housing-dataset', unzip=True)
```

Dataset URL: <https://www.kaggle.com/datasets/altavish/boston-housing-dataset>

3.1 Смотрю, что в данных

```
[3]: df = pd.read_csv('HousingData.csv')
print(df.shape)
df.head()
```

(506, 14)

```
[3]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

3.2 Проверяю типы данных

```
[4]: df.dtypes
```

```
[4]: CRIM      float64
ZN          float64
INDUS       float64
CHAS        float64
NOX         float64
RM          float64
AGE         float64
DIS         float64
RAD         int64
TAX         int64
```

```
PTRATIO    float64
B          float64
LSTAT      float64
MEDV       float64
dtype: object
```

3.3 Проверяю значения категориальных признаков

```
[5]: df.CHAS.unique()
```

```
[5]: array([ 0., nan,  1.])
```

3.4 Проверяю пропуски

```
[6]: df.isna().sum()
```


```
[6]: CRIM      20
     ZN       20
     INDUS   20
     CHAS     20
     NOX      0
     RM       0
     AGE     20
     DIS      0
     RAD      0
     TAX      0
     PTRATIO  0
     B        0
     LSTAT    20
     MEDV     0
     dtype: int64
```

3.5 Заполняю пропуски в численном признаке “CRIM” в соответствии с описанием “CRIM - per capita crime rate by town”

```
[7]: df[df.CRIM == 0]
```


```
[7]: Empty DataFrame
     Columns: [CRIM, ZN, INDUS, CHAS, NOX, RM, AGE, DIS, RAD, TAX, PTRATIO, B, LSTAT, MEDV]
     Index: []
```

```
[8]: df.CRIM.hist(bins=range(90))
     plt.title('Распределение уровня преступности')
     plt.xlabel('Преступлений на душу')
     plt.ylabel('Количество')
     plt.show()
```




```
[9]: df = df.fillna(value={"CRIM": 0})

df.CRIM.hist(bins=range(90))
plt.title('Распределение уровня преступности')
plt.xlabel('Преступлений на душу')
plt.ylabel('Количество')
plt.show()
```




3.6 Заполняю пропуски в численном признаке “ZN” в соответствии с описанием “ZN - proportion of residential land zoned for lots over 25,000 sq.ft.”

```
[10]: df.ZN.hist(bins=range(101))
plt.title('Распределение ZN')
plt.xlabel('Пропорция')
plt.ylabel('Количество')
plt.show()
```




```
[11]: df = df.fillna(value={"ZN": 0})

df.ZN.hist(bins=range(101))
plt.title('Распределение ZN')
plt.xlabel('Пропорция')
plt.ylabel('Количество')
plt.show()
```



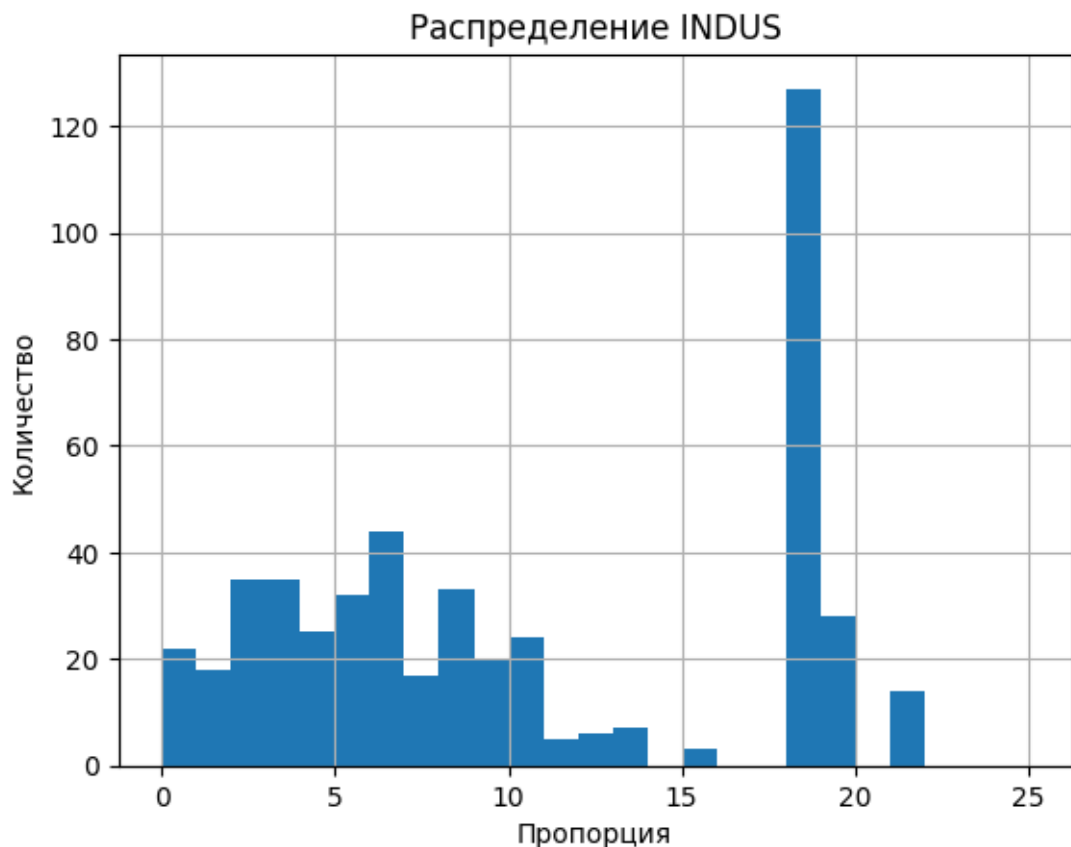
3.7 Заполняю пропуски в численном признаке “INDUS” в соответствии с описанием “INDUS - proportion of non-retail business acres per town.”

```
[12]: df.INDUS.hist(bins=range(26))  
plt.title('Распределение INDUS')  
plt.xlabel('Пропорция')  
plt.ylabel('Количество')  
plt.show()
```



```
[13]: df = df.fillna(value={"INDUS": 0})
```

```
df.INDUS.hist(bins=range(26))  
plt.title('Распределение INDUS')  
plt.xlabel('Пропорция')  
plt.ylabel('Количество')  
plt.show()
```





3.8 Не удаляю пропуски в категориальном признаке “CHAS” в соответствии с описанием “CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise)”

```
[14]: df = df.fillna(value={"CHAS": 2})
```


3.9 Заполняю пропуски в численном признаке “AGE” в соответствии с описанием “AGE - proportion of owner-occupied units built prior to 1940”

```
[15]: df.AGE.hist(bins=range(101))
plt.title('Распределение AGE')
plt.xlabel('Пропорция')
plt.ylabel('Количество')
plt.show()
```



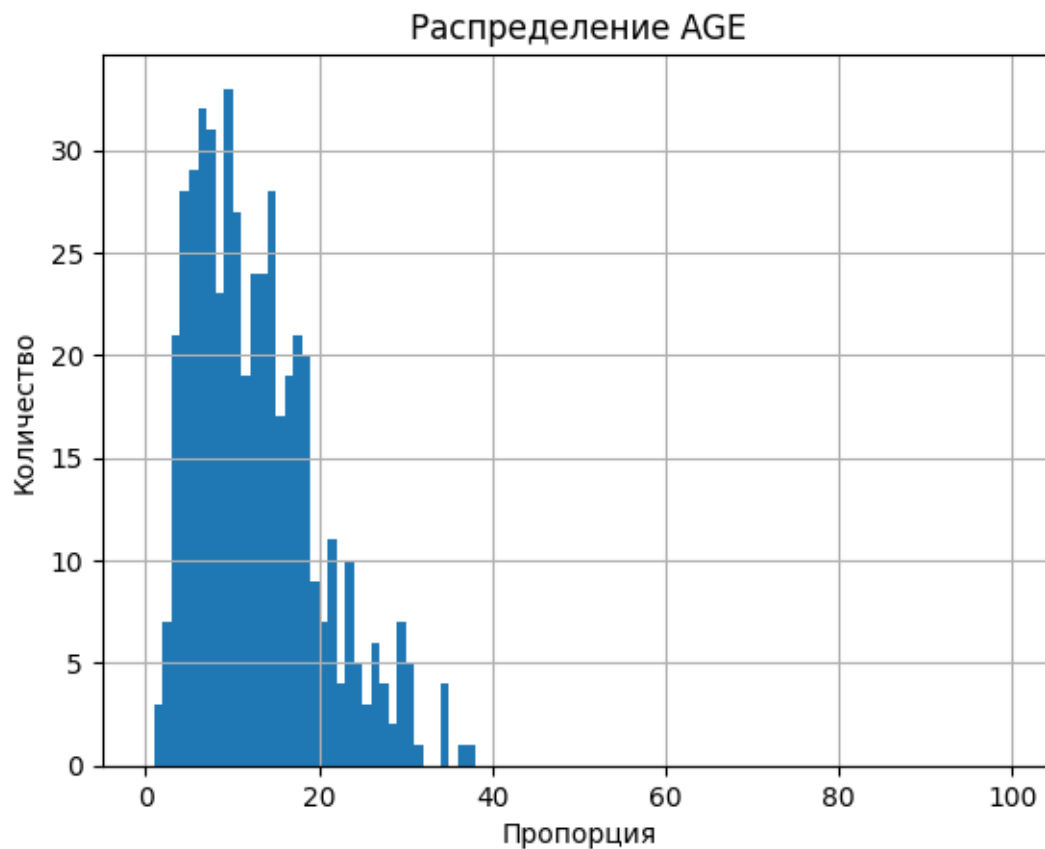
```
[16]: df = df.fillna(value={"AGE": 100})
```

```
df.AGE.hist(bins=range(101))  
plt.title('Распределение AGE')  
plt.xlabel('Пропорция')  
plt.ylabel('Количество')  
plt.show()
```



3.10 Заполняю пропуски в численном признаке “LSTAT” в соответствии с описанием “LSTAT - % lower status of the population”


```
[17]: df.LSTAT.hist(bins=range(101))  
plt.title('Распределение AGE')  
plt.xlabel('Пропорция')  
plt.ylabel('Количество')  
plt.show()
```



```
[18]: med = df.LSTAT.median()
      print(med)
      df = df.fillna(value={"LSTAT": int(med)})

      df.LSTAT.hist(bins=range(101))
      plt.title('Распределение AGE')
      plt.xlabel('Пропорция')
      plt.ylabel('Количество')
      plt.show()
```

11.43



```
[19]: df.isna().sum()
```

```
[19]: CRIM      0
      ZN      0
      INDUS  0
      CHAS   0
      NOX    0
      RM     0
      AGE    0
      DIS    0
      RAD    0
      TAX    0
      PTRATIO 0
      B      0
      LSTAT  0
      MEDV   0
      dtype: int64
```

3.11 Преобразую категориальные признаки (one hot encoding)

```
[20]: for to_enc in ["CHAS"]:  
        one_hot = pd.get_dummies(df[to_enc]).astype(int)  
        del df[to_enc]  
        df = df.join(one_hot)  
df.columns = df.columns.map(str)  
df.head()
```

```
[20]:
```

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.458	7.147	54.2	6.0622	3	222	18.7	

	B	LSTAT	MEDV	0.0	1.0	2.0
0	396.90	4.98	24.0	1	0	0
1	396.90	9.14	21.6	1	0	0
2	392.83	4.03	34.7	1	0	0
3	394.63	2.94	33.4	1	0	0
4	396.90	11.00	36.2	1	0	0

3.12 Провожу разделение на тестовую и обучающую выборки, обучаю и тестирую KNN для предсказания признака MEDV (регрессия), оцениваю с помощью MAE, MSE

```
[21]: def exec_time(start, end):  
        diff_time = end - start  
        m, s = divmod(diff_time, 60)  
        h, m = divmod(m, 60)  
        s, m, h = int(round(s, 0)), int(round(m, 0)), int(round(h, 0))  
        return("{0:02d}:{1:02d}:{2:02d}".format(h, m, s))
```

```
[22]: y = df.MEDV.copy()  
X = df.loc[:, df.columns != "MEDV"].copy()
```

```
[23]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,  
        ↪random_state=42)
```

3.13 RandomForestRegressor

```
[24]: model = RandomForestRegressor(oob_score=True, random_state=42)  
  
start = time.time()  
model.fit(X_train, y_train)  
end = time.time()
```

```

fitTime = exec_time(start, end)

start = time.time()
y_pred = model.predict(X_test)
end = time.time()
testTime = exec_time(start, end)

start = time.time()
y_train_pred = model.predict(X_train)
end = time.time()
trainTime = exec_time(start, end)

testMAE = mean_absolute_error(y_test, y_pred)
trainMAE = mean_absolute_error(y_train, y_train_pred)
testMSE = mean_squared_error(y_test, y_pred)
trainMSE = mean_squared_error(y_train, y_train_pred)
print("Test MAE = %.4f" % testMAE)
print("Train MAE = %.4f" % trainMAE)
print("Test MSE = %.4f" % testMSE)
print("Train MSE = %.4f" % trainMSE)

```

```

Test MAE = 2.1483
Train MAE = 0.9802
Test MSE = 9.7985
Train MSE = 2.5207

```

```

[25]: RandomForestRegressorMAE = pd.DataFrame({
    "Train MAE" : [trainMAE],
    "Test MAE" : [testMAE],
    "Train MSE" : [trainMSE],
    "Test MSE" : [testMSE],
    "Fit time" : [fitTime],
    "Test time on train df" : [trainTime],
    "Test time on test df" : [testTime],
}, index=["RandomForestRegressor"])
RandomForestRegressorMAE

```

```

[25]:
          Train MAE  Test MAE  Train MSE  Test MSE  Fit time  \
RandomForestRegressor    0.980153  2.148323    2.520687  9.798453  00:00:00

          Test time on train df  Test time on test df
RandomForestRegressor          00:00:00          00:00:00

```

3.14 ExtraTreesRegressor

```
[26]: model = ExtraTreesRegressor(bootstrap=True, oob_score=True, random_state=42)

start = time.time()
model.fit(X_train, y_train)
end = time.time()
fitTime = exec_time(start, end)

start = time.time()
y_pred = model.predict(X_test)
end = time.time()
testTime = exec_time(start, end)

start = time.time()
y_train_pred = model.predict(X_train)
end = time.time()
trainTime = exec_time(start, end)

testMAE = mean_absolute_error(y_test, y_pred)
trainMAE = mean_absolute_error(y_train, y_train_pred)
testMSE = mean_squared_error(y_test, y_pred)
trainMSE = mean_squared_error(y_train, y_train_pred)
print("Test MAE = %.4f" % testMAE)
print("Train MAE = %.4f" % trainMAE)
print("Test MSE = %.4f" % testMSE)
print("Train MSE = %.4f" % trainMSE)
```

```
Test MAE = 2.0043
Train MAE = 0.9405
Test MSE = 11.5535
Train MSE = 2.0609
```

```
[27]: ExtraTreesRegressorMAE = pd.DataFrame({
    "Train MAE" : [trainMAE],
    "Test MAE" : [testMAE],
    "Train MSE" : [trainMSE],
    "Test MSE" : [testMSE],
    "Fit time" : [fitTime],
    "Test time on train df" : [trainTime],
    "Test time on test df" : [testTime],
}, index=["ExtraTreesRegressor"])
ExtraTreesRegressorMAE
```

```
[27]:
```

	Train MAE	Test MAE	Train MSE	Test MSE	Fit time	\
ExtraTreesRegressor	0.94051	2.004269	2.060851	11.553487	00:00:00	
						Test time on train df Test time on test df

3.15 AdaBoostRegressor

```
[28]: model = AdaBoostRegressor(n_estimators=100, random_state=42)
```

```
start = time.time()
model.fit(X_train, y_train)
end = time.time()
fitTime = exec_time(start, end)

start = time.time()
y_pred = model.predict(X_test)
end = time.time()
testTime = exec_time(start, end)

start = time.time()
y_train_pred = model.predict(X_train)
end = time.time()
trainTime = exec_time(start, end)

testMAE = mean_absolute_error(y_test, y_pred)
trainMAE = mean_absolute_error(y_train, y_train_pred)
testMSE = mean_squared_error(y_test, y_pred)
trainMSE = mean_squared_error(y_train, y_train_pred)
print("Test MAE = %.4f" % testMAE)
print("Train MAE = %.4f" % trainMAE)
print("Test MSE = %.4f" % testMSE)
print("Train MSE = %.4f" % trainMSE)
```

```
Test MAE = 2.3307
Train MAE = 2.1889
Test MSE = 11.8726
Train MSE = 7.3984
```

```
[29]: AdaBoostRegressorMAE = pd.DataFrame({
    "Train MAE" : [trainMAE],
    "Test MAE" : [testMAE],
    "Train MSE" : [trainMSE],
    "Test MSE" : [testMSE],
    "Fit time" : [fitTime],
    "Test time on train df" : [trainTime],
    "Test time on test df" : [testTime],
}, index=["AdaBoostRegressor"])
AdaBoostRegressorMAE
```

```
[29]:
```

	Train MAE	Test MAE	Train MSE	Test MSE	Fit time	\
AdaBoostRegressor	2.188942	2.330673	7.398446	11.872565	00:00:00	

	Test time on train df	Test time on test df
AdaBoostRegressor	00:00:00	00:00:00

3.16 GradientBoostingRegressor

```
[30]: model = GradientBoostingRegressor(random_state=42)

start = time.time()
model.fit(X_train, y_train)
end = time.time()
fitTime = exec_time(start, end)

start = time.time()
y_pred = model.predict(X_test)
end = time.time()
testTime = exec_time(start, end)

start = time.time()
y_train_pred = model.predict(X_train)
end = time.time()
trainTime = exec_time(start, end)

testMAE = mean_absolute_error(y_test, y_pred)
trainMAE = mean_absolute_error(y_train, y_train_pred)
testMSE = mean_squared_error(y_test, y_pred)
trainMSE = mean_squared_error(y_train, y_train_pred)
print("Test MAE = %.4f" % testMAE)
print("Train MAE = %.4f" % trainMAE)
print("Test MSE = %.4f" % testMSE)
print("Train MSE = %.4f" % trainMSE)
```

```
Test MAE = 2.0407
Train MAE = 1.0978
Test MSE = 8.4135
Train MSE = 1.9931
```

```
[31]: GradientBoostingRegressorMAE = pd.DataFrame({
    "Train MAE" : [trainMAE],
    "Test MAE" : [testMAE],
    "Train MSE" : [trainMSE],
    "Test MSE" : [testMSE],
    "Fit time" : [fitTime],
    "Test time on train df" : [trainTime],
    "Test time on test df" : [testTime],
}, index=["GradientBoostingRegressor"])
```

```
GradientBoostingRegressorMAE
```

```
[31]:
```

	Train MAE	Test MAE	Train MSE	Test MSE	Fit time \
GradientBoostingRegressor	1.097762	2.04074	1.99313	8.413547	00:00:00

	Test time on train df	Test time on test df
GradientBoostingRegressor	00:00:00	00:00:00

3.17 Провожу сравнение

```
[32]: AllMAE = pd.concat([RandomForestRegressorMAE, ExtraTreesRegressorMAE,
↳AdaBoostRegressorMAE, GradientBoostingRegressorMAE])
AllMAE.sort_values(by=["Test MSE"])
```

```
[32]:
```

	Train MAE	Test MAE	Train MSE	Test MSE \
GradientBoostingRegressor	1.097762	2.040740	1.993130	8.413547
RandomForestRegressor	0.980153	2.148323	2.520687	9.798453
ExtraTreesRegressor	0.940510	2.004269	2.060851	11.553487
AdaBoostRegressor	2.188942	2.330673	7.398446	11.872565

	Fit time	Test time on train df	Test time on test df
GradientBoostingRegressor	00:00:00	00:00:00	00:00:00
RandomForestRegressor	00:00:00	00:00:00	00:00:00
ExtraTreesRegressor	00:00:00	00:00:00	00:00:00
AdaBoostRegressor	00:00:00	00:00:00	00:00:00